

Context Engineering for Claude Code

The INAV-Claude Project

Sensei

Making Claude Code work better through context management

The Problem

Common Problems Without Context Engineering

- Forgot to check lock files
- Used `make` instead of inav-builder agent
- Skipped testing before PR
- Didn't run code review
- Pushed to master instead of feature branch
- Missing commit message details
- Loaded 100k lines but missed the critical 100

 **Claude isn't bad - the process isn't structured!**

Why This Happens

```
| 100k Lines of Code |  
| + 5k Lines of Documentation |  
| + Build Instructions |  
| + Testing Guides |  
| + Project Tracking |  
| + Architecture Docs |  
| = 110k+ Lines Total |
```



```
| Claude's Context | ← Information overload  
| [saturated] |
```



```
| Important info |  
| gets buried! |
```

Solution Part 1 - Role Separation



MANAGER

Context Size: ~1,200 lines

Focus: Planning & Coordination

Loads: Project tracking,
INDEX.md

Doesn't Load: Build instructions



RELEASE MANAGER

Context Size: ~1,000 lines

Focus: Builds & Artifacts

Loads: Release workflow,
changelogs

Doesn't Load: Implementation



DEVELOPER

Each role sees ONLY what

Solution Part 2 - Communication System

```
# Task: Fix Terrain Data Not Loading
```

```
## Priority: HIGH
```

```
## Problem
```

User reports: "terrain data doesn't load" in Mission Control.

```
## Success Criteria           ← Clear goals
```

- [] Root cause identified
- [] Terrain data loads successfully
- [] PR created with tests

```
## Available Resources        ← What you have
```

- Chrome DevTools MCP available
- test-engineer agent for UI testing

Solution Part 3 - 12-Step Workflow + JIT Docs

Developer 12-Step Process: Documentation Loads:

1. Check inbox
2. Read task [Task file: 80 lines]
3. Create git branch → [CRITICAL-BEFORE-CODE: 104 lines]
 - Check lock files
4. Reproduce bug (fails) —
 - Use agents, not direct commands

|
5. Implement fix | [CRITICAL-BEFORE-TEST: 113 lines]

→ • Test philosophy
6. Compile code | • Edge cases

|
7. Verify fix (passes) —

Solution Part 4 - Specialized Agents



inav-builder

Agent File: 282 lines
Represents: ~3,000 lines

Knows:

- CMake build system
- ARM cross-compilation
- Linker compatibility

Doesn't Know:

- Mission planning
- MSP protocol



test-engineer

Agent File: 492 lines
Represents: ~2,500 lines

Knows:

- Chrome DevTools Protocol
- UI testing strategies
- SITL simulator usage

Doesn't Know:

- Build systems
- Project management

How It All Fits Together

USER: "Fix GPS bug" |



Developer Role | [Focused Context]
Loads: 2,500 lines |



/start-task skill | [Reusable Workflow]
• Check locks |
• Create branch |



Execute 12-Step Workflow |



Hooks - Automatic Enforcement

WITHOUT HOOKS

Claude: [runs `make SITL`]

Result:  Build fails

- Wrong directory
- Missing flags

User: "Use the build script"

Claude: "Oh, sorry!"

[Next task...]

Claude: [runs `make SITL` again]

User: 

WITH HOOKS

Claude: [tries `make SITL`]

Hook:  **INTERCEPTED**

Hook: "Use inav-builder agent"

Claude: [uses inav-builder]

Result:  Build succeeds

[Next task...]

Claude: [starts typing "make"]

Hook:  **INTERCEPTED**

Real Example - Fix Terrain Data Loading

Screenshots advancing every 10 seconds (1:40 total)

0:00 → Problem: Terrain chart not displaying

0:10 → Task assigned: 80-line structured file

0:20 → Developer starts, guide loads

0:30 → test-engineer investigates with DevTools

0:40 → Root cause: `plotElevation()` commented out

0:50 → Fix: Chart.js v4 integration

1:00 → inav-code-review checks quality

1:10 → Success: Chart displays correctly

1:20 → PR #2518 created

1:30 → Context comparison

Results from Real-World Use



PRODUCTIVITY & CONSISTENCY

Projects Completed: **78 in last 2 months**

Same-Day Completions: 15+

- fix-terrain-data: 4 hours
- fix-blackbox: 1-word fix
- fix-climb-rate: 1 operator

Process Consistency: **VERY HIGH**

- ✓ Projects follow 12-step workflow
- ✓ Testing before PR



CONTEXT EFFICIENCY

Codebase: 150,000 lines

Typical task loads:

- ~1,500 lines (with system)
- ~10-15k lines (without)

Result:

- Faster responses
- Fewer mistakes
- Better guideline adherence
- Professional dev process

5 Principles You Can Use

- ✓ 1. STRUCTURE BY ROLE AND PHASE
- ✓ 2. LOAD DOCS JUST-IN-TIME
- ✓ 3. USE SPECIALIZED AGENTS
- ✓ 4. ENFORCE WITH HOOKS
- ✓ 5. CLEAR COMMUNICATION BOUNDARIES

These patterns work for ANY large codebase

Self-Improvement

System Creates Its Own Tools

create-agent builds new agents!

User: "We keep looking up MSP messages"

Claude:

1. Researches MSP docs
2. Designs msp-expert agent
3. Writes agent (271 lines)
4. Updates README

Our project improves its own

Lessons Learned

Guides have self-documentation:

Lessons Learned

- ****Lock file format****: Include timestamp for debugging
- ****inav-architecture first****: Use before Grep - saves 10min
- ****SITL directory****: Use `build_sitl/` *not* `build/`

System documents itself as it learns

Adapt for Your Project



**THE WORKFLOW IS
UNIVERSAL**

**12 steps work for most
development projects**

Firmware, web apps, data science...

The STRUCTURE is reusable

- ✓ Role separation
- ✓ JIT documentation
- ✓ Agent pattern
- ✓ Hook enforcement

The CONTENT needs customization

- ✗ Your build commands



GETTING STARTED

Week 1: Role separation

Week 2: JIT guides

Week 3: First agent

Week 4: Add hooks

Examples:

Python/Django, React/TypeScript, Rust

Thank You!

Questions?

Repository: github.com/sensei-hacker/inav-claude

Contact: sensei-hacker on GitHub

*Context engineering turns Claude from a smart assistant
into a reliable, professional development team member*

Backup: File Structure

```
inavflight/
└── .claude/ # Claude Code configuration
    ├── settings.json # Hooks, permissions
    ├── agents/ # 10 agents, 3,301 lines
    │   ├── inav-builder.md
    │   ├── test-engineer.md
    │   └── msp-expert.md
    ├── skills/ # 31 reusable workflows
    └── hooks/ # Enforcement scripts

    └── claude/ # Role-specific workspaces
        ├── manager/
        │   └── README.md # 1,200 lines of context
        ├── developer/
        │   ├── README.md # 2,500 lines of context
        │   └── guides/
        │       ├── CRITICAL-BEFORE-CODE.md (104 lines)
        │       ├── CRITICAL-BEFORE-TEST.md (113 lines)
        │       └── CRITICAL-BEFORE-COMMIT.md (105 lines)
```